

Utilizing Parallel Workers: LLVM's Vectorization Plan

Jonas Fritsch

Technical University of Munich

jonas.fritsch@tum.de

Abstract

abstract

- (potential security vulnerabilities [11])

1 Introduction

1. Modern CPUs (SIMD-Registers, SIMD ISAs)
2. SIMD performance gains
3. Auto-Vectorization to help programmers utilize SIMD functionality out of the box (explicit vectorization as time consuming especially when supporting multiple architectures)
4. LLVM as popular, platform-agnostic compiler framework [14]
5. auto-vectorization complexity → wish for well designed, interoperable system (VPlan) → transition from Loop Vectorizer + SLP Vectorizer to Vectorization Plan → ongoing 8 year long effort started by Intel

2 Background

1. Vectorization explanation (with Code-Examples)
 - Loop-Level Vectorization (Inner / Outer-Loops [4, 8, 18, 21, 22])
 - Function Vectorization [21]
 - Superword-Level Parallelism Vectorization [13] (maybe already [5])
2. Vectorization Constraints [2, 9] (reuse code-examples from above)
 - not always legal / applicable, e.g.: backwards data-dependencies, pointer aliasing, platform backwards-compatibility / support, control flow complexity (function calls, register pressure), FP inaccuracies (`-ffast-math`)
 - possible performance loss, e.g.: costly conversions (horizontal aggregations, integer division) resulting in slower overall code on average
 - larger code size, e.g.: need for scalar loop-tail / epilogue

3 LLVM's Vectorization Plan

Overview [4, 7, 20, 21, 24, 27, 28, 30]

1. History
 - Short Summary how it was before (LoopVectorizer)
 - Initial refactoring proposition by Intel
 - Describe rough final finished VPlan architecture.
2. State of VPlan
 - Rough algorithm description (interplay of legacy LV, legacy cost model and new VPlan) (with light code references?)
 - VPlan structure (HCFG, `VPRegionBlock`, recipes, etc.) and modelling of IR explained via example similar to [30] (What is good example? Too simple → doesn't cover enough; Too complex → doesn't fit into paper.)
 - (Outer Loop Vectorization?)
3. Future of VPlan (what is currently being worked on, what is planned next, what has not been planned)

4 Related Work

1. Region Vectorizer (RV) [16, 24]
2. (Other auto-vectorization algorithms / possible improvements [2, 12, 17, 19, 23])
3. (Auto-Vectorization in GCC [10, 26])
4. (Comparison with GCC [6, 15])
5. (Comparison/Outlook auto-vectorization and explicit vectorization [1, 3, 6])
6. (Polyhedral compilation techniques [29])
7. (LLM-based Vectorization [25])

5 Summary and Future Work

VPlan Summary

References

- [1] Neil Adit and Adrian Sampson. 2022. Performance Left on the Table: An Evaluation of Compiler Autovectorization for RISC-V. *IEEE Micro* 42, 5 (2022), 41–48. <https://doi.org/10.1109/MM.2022.3184867>
- [2] Anupama Rasale Ashutosh Nema. 2023. Improving Vectorization for Loops with Control Flow. <https://www.youtube.com/watch?v=mKG0NmGtpbE> Accessed: 2024-11-04.
- [3] Lawrence Benson, Richard Ebeling, and Tilmann Rabl. 2023. Evaluating SIMD Compiler-Intrinsics for Database Systems. In *Joint Proceedings of Workshops at the 49th International Conference on Very Large Data Bases (VLDB 2023), Vancouver, Canada, August 28 - September 1, 2023 (CEUR Workshop Proceedings)*, Vol. 3462. CEUR-WS.org. <https://ceur-ws.org/Vol-3462/ADMS5.pdf>
- [4] Diego Caballero and Intel Vectorizer Team. 2018. Extending LoopVectorize to Support Outer Loop Vectorization Using VPlan. <https://llvm.org/devmtg/2018-04/slides/Caballero-Extending%20LoopVectorize%20to%20Support%20Outer%20Loop%20Vectorization%20Using%20VPlan.pdf> Accessed: 2024-11-18.
- [5] Yishen Chen, Charith Mendis, and Saman Amarasinghe. 2022. All you need is superword-level parallelism: systematic control-flow vectorization with SLP. In *Proceedings of the 43rd ACM SIGPLAN International Conference on Programming Language Design and Implementation (PLDI 2022)*. Association for Computing Machinery, New York, NY, USA, 301–315. <https://doi.org/10.1145/3519939.3523701>
- [6] Jing Ge Feng, Ye Ping He, Qiu Ming Tao, and Fazli Wahid. 2021. Evaluation of Compilers’ Capability of Automatic Vectorization Based on Source Code Analysis. *Sci. Program.* 2021 (Nov. 2021), 15. <https://doi.org/10.1155/2021/3264624>
- [7] Florian Hahn and Ayal Zaks. 2023. VPlan: Status Update and Future Roadmap. <https://llvm.org/devmtg/2023-10/slides/techtalks/Hahn-VPlan-StatusUpdateAndRoadmap.pdf> Accessed: 2024-11-19.
- [8] Intel. 2019. Outer Loop Vectorization. <https://www.intel.com/content/www/us/en/developer/articles/technical/outer-loop-vectorization.html> Accessed: 2024-11-06.
- [9] Intel. 2019. Vectorization Essentials. <https://www.intel.com/content/www/us/en/developer/articles/technical/vectorization-essential.html> Accessed: 2024-11-11.
- [10] Jakub Jelínek. 2023. Vectorization optimization in GCC. <https://developers.redhat.com/articles/2023/12/08/vectorization-optimization-gcc> Accessed: 2024-11-04.
- [11] Sayinath Karuppanan and Samira Mirbagher Ajor-paz. 2023. An Attack on The Speculative Vectorization: Leakage from Higher Dimensional Speculation. *arXiv e-prints*, Article arXiv:2302.01131 (Feb. 2023), 15 pages. <https://doi.org/10.48550/arXiv.2302.01131> arXiv:cs.CR/2302.01131
- [12] Universität-Saarland Compiler Design Lab. 2024. RV: A Unified Region Vectorizer for LLVM. <https://github.com/cdl-saarland/rv> Accessed: 2024-11-11.
- [13] Samuel Larsen and Saman Amarasinghe. 2000. Exploiting superword level parallelism with multimedia instruction sets. In *Proceedings of the ACM SIGPLAN 2000 Conference on Programming Language Design and Implementation (PLDI ’00)*. Association for Computing Machinery, New York, NY, USA, 145–156. <https://doi.org/10.1145/349299.349320>
- [14] Chris Lattner and Vikram Adve. 2004. LLVM: A Compilation Framework for Lifelong Program Analysis & Transformation. In *Proceedings of the International Symposium on Code Generation and Optimization: Feedback-Directed and Runtime Optimization (CGO ’04)*. IEEE Computer Society, USA, 75.
- [15] Klara Modin. 2024. A comparison of auto-vectorization performance between GCC and LLVM for the RISC-V vector extension. (10 2024). <https://urn.kb.se/resolve?urn=urn:nbn:se:kth:diva-354873>
- [16] Simon Moll and Sebastian Hack. 2018. Partial control-flow linearization. In *Proceedings of the 39th ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI 2018)*. Association for Computing Machinery, New York, NY, USA, 543–556. <https://doi.org/10.1145/3192366.3192413>
- [17] Simon Moll, Shrey Sharma, Matthias Kurtenacker, and Sebastian Hack. 2019. Multi-dimensional Vectorization in LLVM. In *Proceedings of the 5th Workshop on Programming Models for SIMD/Vector Processing (WP-MVP’19)*. Association for Computing Machinery, New York, NY, USA, Article 3, 8 pages. <https://doi.org/10.1145/3303117.3306172>
- [18] Dorit Nuzman and Ayal Zaks. 2008. Outer-loop vectorization: revisited for short SIMD architectures. In *Proceedings of the 17th International Conference on Parallel Architectures and Compilation Techniques (PACT ’08)*. Association for Computing Machinery, New York, NY, USA, 2–11. <https://doi.org/10.1145/1454115.1454119>
- [19] Rouzbeh Paktinatkeleshteri, João P. L de Carvalho, Ehsan Amiri, and J. Nelson Amaral. 2023. Efficient Auto-Vectorization for Control-flow Dependent Loops through Data Permutation. In *Proceedings of the 33rd Annual International Conference on Computer Science and Software Engineering (CASCON ’23)*. IBM Corp., USA, 74–83.
- [20] Gil Rapaport and Ayal Zaks. 2017. Introducing VPlan to the Loop Vectorizer. https://llvm.org/devmtg/2017-03/assets/slides/introducing_vplan_to_the_loop_vectorizer.pdf Accessed: 2024-11-16.
- [21] Hideki Saito and Intel Vectorizer Team. 2016. Extending LoopVectorizer towards supporting OpenMP4.5 SIMD and outer loop auto-vectorization. <https://llvm.org/devmtg/2016-11/Slides/Saito-NextLevelLLVMLoopVectorizer.pdf> Accessed: 2024-11-18.

Utilizing Parallel Workers: LLVM's Vectorization Plan

- [22] Nadathur Satish, Changkyu Kim, Jatin Chhugani, Hideki Saito, Rakesh Krishnaiyer, Mikhail Smelyanskiy, Milind Girkar, and Pradeep Dubey. 2012. Can traditional programming bridge the Ninja performance gap for parallel computing applications?. In *Proceedings of the 39th Annual International Symposium on Computer Architecture (ISCA '12)*. IEEE Computer Society, USA, 440–451.
- [23] Jaewook Shin. 2007. Introducing Control Flow into Vectorized Code. In *16th International Conference on Parallel Architecture and Compilation Techniques (PACT 2007)*. 280–291. <https://doi.org/10.1109/PACT.2007.4336219>
- [24] Sebastian Hack Simon Moll. 2017. VPlan + Rv: A Proposal. <https://www.youtube.com/watch?v=sMEphbFukw> Accessed: 2024-11-11.
- [25] Jubi Taneja, Avery Laird, Cong Yan, Madan Musuvathi, and Shuvendu K. Lahiri. 2024. LLM-Vectorizer: LLM-based Verified Loop Vectorizer. *arXiv:cs.SE/2406.04693* <https://arxiv.org/abs/2406.04693>
- [26] GCC Team. 2023. Auto-vectorization in GCC. <https://gcc.gnu.org/projects/tree-ssa/vectorization.html> Accessed: 2024-11-04.
- [27] LLVM Team. 2021. Auto-Vectorization in LLVM. <https://llvm.org/docs/Vectorizers.html> Accessed: 2024-11-18.
- [28] LLVM Team. 2024. Vectorization Plan. <https://llvm.org/docs/VectorizationPlan.html> Accessed: 2024-11-18.
- [29] Polly Team. 2017. Polly - LLVM Framework for High-Level Loop and Data-Locality Optimizations. <https://polly.llvm.org/index.html> Accessed: 2024-11-11.
- [30] Ayal Zaks and Gil Rapaport. 2017. Vectorizing Loops with VPlan – Current State and Next Steps. <https://llvm.org/devmtg/2017-10/slides/Zaks-Vectorizing%20Loops%20with%20VPlan.pdf> Accessed: 2024-11-16.