

Utilizing Parallel Workers: LLVM's Vectorization Plan

Jonas Fritsch

Technical University of Munich

jonas.fritsch@tum.de

Abstract

Lorem ipsum.

1 Introduction

1. Modern CPUs (SIMD-Registers, SIMD ISAs)
2. SIMD performance gains
3. Auto-Vectorization to help programmers better utilize SIMD functionality out of the box
4. LLVM as popular, platform-agnostic compiler framework
5. auto-vectorization complexity → wish for well designed, interoperable system (VPlan) → ongoing 8 year long effort started by Intel
6. rough paper outline

2 Background

1. Vectorization explanation (with Code-Examples)
 - Loop-Level Vectorization (Inner / Outer-Loops [6, 10, 18, 22, 23])
 - Function Vectorization [22]
 - Superword-Level Parallelism Vectorization [15] (maybe already [7])
2. Vectorization Drawbacks [2, 11] (reuse code-examples from above)
 - not always legal / applicable, e.g.: backwards data-dependencies, pointer aliasing, platform backwards-compatibility / support, control flow complexity (function calls, register pressure), FP inaccuracies (`-ffast-math`)
 - possible performance loss, e.g.: costly conversions (horizontal aggregations, integer division) resulting in slower overall code on average
 - larger code size, e.g.: need for scalar loop-tail / epilogue
 - (potential security vulnerabilities [13])
3. OpenMP (very short)
4. Auto-Vectorization in Compilers (short)

5. LLVM (very brief introduction) [8]
6. LLVM 'recent' transition from Loop Vectorizer + SLP Vectorizer to Vectorization Plan (ongoing 8 year long refactoring and improvement effort started by Intel)

3 LLVM's Vectorization Plan

Overview [3, 4, 6, 21, 22, 25, 28, 29]

1. Phase 1: Legality
2. Phase 2: Planning (Vectorization Plans)
 - a. Building initial VPlan based on Phase 1 Constraints (Generate one or more up-to-date Code Examples for VPlan similar to [4, 6, 21]) → Explain different components (e.g.: `VPBasicBlock`, `VPRegionBlock`, etc.) and correlation between generated VPlan structure, original scalar code. Show some general optimizations.
 - b. Modelling Cost. Not yet part of VPlan. (related work? already adopted? [20])
 - c. Optimizing/Transforming VPlans [3]
3. Phase 3: Execution - Materialize best VPlan

4 Related Work

1. Other auto-vectorization algorithms / possible improvements [2, 14, 17, 19, 24]
Some of those should already be (partially) implemented in LLVM
2. Auto-Vectorization in GCC [12, 27]
3. Comparison with GCC [9, 16]
4. (Polyhedral compilation techniques [30])
5. (LLM-based Vectorization [26])

5 Summary and Future Work

1. VPlan Summary [3, 29]
2. VPlan Future Roadmap [3, 29] (e.g.: combining outer-/inner-loop paths, full function vectorization, etc.)
3. (Comparison/Outlook auto-vectorization and explicit vectorization [1, 5, 9])

References

- [1] Neil Adit and Adrian Sampson. 2022. Performance Left on the Table: An Evaluation of Compiler Autovectorization for RISC-V. *IEEE Micro* 42, 5 (2022), 41–48. <https://doi.org/10.1109/MM.2022.3184867>
- [2] Anupama Rasale Ashutosh Nema. 2023. Improving Vectorization for Loops with Control Flow. <https://www.youtube.com/watch?v=mKG0NmGtpbE> Accessed: 2024-11-04.
- [3] Florian Hahn Ayal Zaks. 2023. VPlan: Status Update and Future Roadmap. <https://www.youtube.com/watch?v=SzGP4PgMuLE> Accessed: 2024-11-02.
- [4] Gil Rapaport Ayal Zaks. 2017. Vectorizing Loops with VPlan – Current State and Next Steps. <https://www.youtube.com/watch?v=BjBSJFzYDVk> Accessed: 2024-11-04.
- [5] Lawrence Benson, Richard Ebeling, and Tilmann Rabl. 2023. Evaluating SIMD Compiler-Intrinsics for Database Systems. In *Joint Proceedings of Workshops at the 49th International Conference on Very Large Data Bases (VLDB 2023), Vancouver, Canada, August 28 - September 1, 2023 (CEUR Workshop Proceedings)*, Vol. 3462. CEUR-WS.org. <https://ceur-ws.org/Vol-3462/ADMS5.pdf>
- [6] Diego Caballero and Vectorizer Team Intel Corporation. 2018. Extending LoopVectorizer to Support Outer Loop Vectorization Using VPlan. <https://www.youtube.com/watch?v=z6NeHLRNvok> Accessed: 2024-11-04.
- [7] Yishen Chen, Charith Mendis, and Saman Amarasinghe. 2022. All you need is superword-level parallelism: systematic control-flow vectorization with SLP. In *Proceedings of the 43rd ACM SIGPLAN International Conference on Programming Language Design and Implementation (PLDI 2022)*. Association for Computing Machinery, New York, NY, USA, 301–315. <https://doi.org/10.1145/3519939.3523701>
- [8] Vikram Adve Chris Lattner. 2004. LLVM: A Compilation Framework for Lifelong Program Analysis & Transformation. In *Proceedings of the International Symposium on Code Generation and Optimization: Feedback-Directed and Runtime Optimization (CGO '04)*. IEEE Computer Society, USA, 75.
- [9] Jing Ge Feng, Ye Ping He, Qiu Ming Tao, and Fazli Wahid. 2021. Evaluation of Compilers' Capability of Automatic Vectorization Based on Source Code Analysis. *Sci. Program.* 2021 (Nov. 2021), 15. <https://doi.org/10.1155/2021/3264624>
- [10] Intel. 2019. Outer Loop Vectorization. <https://www.intel.com/content/www/us/en/developer/articles/technical/outer-loop-vectorization.html> Accessed: 2024-11-06.
- [11] Intel. 2019. Vectorization Essentials. <https://www.intel.com/content/www/us/en/developer/articles/technical/vectorization-essential.html> Accessed: 2024-11-11.
- [12] Jakub Jelínek. 2023. Vectorization optimization in GCC. <https://developers.redhat.com/articles/2023/12/08/vectorization-optimization-gcc> Accessed: 2024-11-04.
- [13] Sayinath Karuppanan and Samira Mirbagher Ajor-paz. 2023. An Attack on The Speculative Vectorization: Leakage from Higher Dimensional Speculation. *arXiv e-prints*, Article arXiv:2302.01131 (Feb. 2023), 15 pages. <https://doi.org/10.48550/arXiv.2302.01131>
- [14] Universität-Saarland Compiler Design Lab. 2024. RV: A Unified Region Vectorizer for LLVM. <https://github.com/cdl-saarland/rv> Accessed: 2024-11-11.
- [15] Samuel Larsen and Saman Amarasinghe. 2000. Exploiting superword level parallelism with multimedia instruction sets. In *Proceedings of the ACM SIGPLAN 2000 Conference on Programming Language Design and Implementation (PLDI '00)*. Association for Computing Machinery, New York, NY, USA, 145–156. <https://doi.org/10.1145/349299.349320>
- [16] Klara Modin. 2024. A comparison of auto-vectorization performance between GCC and LLVM for the RISC-V vector extension. (10 2024). <https://urn.kb.se/resolve?urn=urn:nbn:se:kth:diva-354873>
- [17] Simon Moll, Shrey Sharma, Matthias Kurtenacker, and Sebastian Hack. 2019. Multi-dimensional Vectorization in LLVM. In *Proceedings of the 5th Workshop on Programming Models for SIMD/Vector Processing (WPMVP'19)*. Association for Computing Machinery, New York, NY, USA, Article 3, 8 pages. <https://doi.org/10.1145/3303117.3306172>
- [18] Dorit Nuzman and Ayal Zaks. 2008. Outer-loop vectorization: revisited for short SIMD architectures. In *Proceedings of the 17th International Conference on Parallel Architectures and Compilation Techniques (PACT '08)*. Association for Computing Machinery, New York, NY, USA, 2–11. <https://doi.org/10.1145/1454115.1454119>
- [19] Rouzbeh Paktinatkeleshteri, João P. L de Carvalho, Ehsan Amiri, and J. Nelson Amaral. 2023. Efficient Auto-Vectorization for Control-flow Dependent Loops through Data Permutation. In *Proceedings of the 33rd Annual International Conference on Computer Science and Software Engineering (CASCON '23)*. IBM Corp., USA, 74–83.
- [20] Angela Pohl, Biagio Cosenza, and Ben Juurlink. 2020. Vectorization cost modeling for NEON, AVX and SVE. *Performance Evaluation* 140-141 (2020), 102106. <https://doi.org/10.1016/j.peva.2020.102106>
- [21] Gil Rapaport. 2017. Introducing VPlan to the Loop Vectorizer. <https://www.youtube.com/watch?v=lqzJR56tb7Y> Accessed: 2024-11-06.
- [22] Hideki Saito. 2016. Extending LoopVectorizer towards supporting OpenMP4.5 SIMD and outer loop auto-vectorization. <https://www.youtube.com/watch?v=XXAvdUwO7kQ> Accessed: 2024-11-06.
- [23] Nadathur Satish, Changkyu Kim, Jatin Chhugani, Hideki Saito, Rakesh Krishnaiyer, Mikhail Smelyanskiy, Milind Girkar, and Pradeep Dubey. 2012. Can traditional programming bridge the Ninja performance gap for parallel computing applications?. In *Proceedings of the 39th Annual International Symposium on Computer Architecture (ISCA '12)*. IEEE Computer Society, USA,

Utilizing Parallel Workers: LLVM's Vectorization Plan

440–451.

- [24] Jaewook Shin. 2007. Introducing Control Flow into Vectorized Code. In *16th International Conference on Parallel Architecture and Compilation Techniques (PACT 2007)*. 280–291. <https://doi.org/10.1109/PACT.2007.4336219>
- [25] Sebastian Hack Simon Moll. 2017. VPlan + Rv: A Proposal. <https://www.youtube.com/watch?v=svMEphbFukw> Accessed: 2024-11-11.
- [26] Jubi Taneja, Avery Laird, Cong Yan, Madan Musuvathi, and Shuvendu K. Lahiri. 2024. LLM-Vectorizer: LLM-based Verified Loop Vectorizer. arXiv:cs.SE/2406.04693 <https://arxiv.org/abs/2406.04693>
- [27] GCC Team. 2023. Auto-vectorization in GCC. <https://gcc.gnu.org/projects/tree-ssa/vectorization.html> Accessed: 2024-11-04.
- [28] LLVM Team. 2021. Auto-Vectorization in LLVM. <https://llvm.org/docs/Vectorizers.html> Accessed: 2024-11-04.
- [29] LLVM Team. 2024. Vectorization Plan. <https://llvm.org/docs/VectorizationPlan.html> Accessed: 2024-11-04.
- [30] Polly Team. 2017. Polly - LLVM Framework for High-Level Loop and Data-Locality Optimizations. <https://polly.llvm.org/index.html> Accessed: 2024-11-11.